

Active Learning

Laura, Peter, Simon

We assume a prior $\pi(w)$ where \mathbf{w} is a vector of parameters that describe our model. As the model we use a sigmoid function

$$\sigma(w_0, w_1, x) = \frac{1}{1 + \exp[-(w_0 + w_1 x)]} \quad (1)$$

NB: This might not be the best way to write the sigmoid. We collect data n data points by presenting a stimulus $x \in \mathbb{R}$ and observing a binary response $y \in \{0, 1\}$:

$$D = \{(x_1, y_1), \dots, (x_n, y_n)\} \equiv (X, Y) \quad (2)$$

The likelihood of the parameters \mathbf{w} given the data D is given by

$$P(Y | \mathbf{w}, X) = \prod_{i=1}^n P(y_i | \mathbf{w}, x_i) \quad (3)$$

$$= \prod_{i=1}^n \sigma(w_0 + w_1 x_i)^{y_i} (1 - \sigma(w_0 + w_1 x_i))^{1-y_i} \quad (4)$$

NB: The stimuli x are not considered part of the data, because we have control over it The posterior probability of the parameters \mathbf{w} is

$$P(\mathbf{w} | X, Y) = \frac{P(Y | \mathbf{w}, X) \pi(\mathbf{w})}{P(Y | X)} \quad (5)$$

The denominator, i.e. the marginal likelihood, is computed by taking the integral over all hypotheses:

$$\int P(Y \mid \mathbf{w}', X) \pi(\mathbf{w}') d\mathbf{w}' = \iint P(Y \mid w_0, w_1, X) \pi(w_0, w_1) dw'_0 dw'_1 \quad (6)$$

The goal is to get a posterior $P(\mathbf{w} \mid X, Y)$ that is of low uncertainty. We use entropy as a measure of the current uncertainty of our estimation of \mathbf{w} . By *current* we mean that we use the data we have discovered in the n steps until now. To make this clear we write X_n, Y_n instead of X, Y :

$$H(\mathbf{w} \mid X_n Y_n) = - \int P(\mathbf{w}' \mid X_n, Y_n) \log[P(\mathbf{w}' \mid X_n, Y_n)] d\mathbf{w}'. \quad (7)$$

In principle we would now like to choose our next stimulus x_{n+1} such that it minimizes the resulting entropy $H(\mathbf{w} \mid X_n, Y_n, x_{n+1}, y_{n+1})$, but we do not know what y_{n+1} is. So we want to find the x_{n+1} that minimizes the mean:

$$K(x_{n+1}) = H(\mathbf{w} \mid X_n, Y_n, x_{n+1}, y_{n+1} = 0) P(y_{n+1} = 0 \mid X_n, Y_n, x_{n+1}) \quad (8)$$

$$+ H(\mathbf{w} \mid X_n, Y_n, x_{n+1}, y_{n+1} = 1) P(y_{n+1} = 1 \mid X_n, Y_n, x_{n+1}). \quad (9)$$

Here $P(y_{n+1} = 0/1 \mid X_n, Y_n, x_{n+1})$ is called the *predictive distribution*. Determining them again requires an integral over the hypotheses:

$$P(y_{n+1} = 0/1 \mid X_n, Y_n, x_{n+1}) = \int P(y_{n+1} = 0/1 \mid \mathbf{w}', x_{n+1}) P(\mathbf{w}' \mid X_n, Y_n) d\mathbf{w}' \quad (10)$$

⁴ For